

Introduction to Artificial Intelligence

Lecture 0: Artificial Intelligence

ChatGPT PLUS

Help me pick
a gift for my dad who loves fishing

Brainstorm edge cases
for a function with birthdate as input, horoscope as ou...

Make up a story
about Sharky, a tooth-brushing shark superhero

Create a personal webpage for me
after asking me three questions

+ Send a message



10:27 🏠 ⏱ 0.00 KB/S VOLTE 86 % 🔋

X




⋮

08/09, 10:26 | 45

So tell me, what would you recommend for a
1-day trip to |

One simple idea:

Guess the next word



In the 1960s, Armstrong ____

In the 1960s, Armstrong performed ____

In the 1960s, Armstrong performed a moonwalk ____

In the 1960s, Armstrong performed a moonwalk on the ____

In the 1960s, Armstrong performed a moonwalk on the lunar ____

In the 1960s, Armstrong performed a moonwalk on the lunar surface and said ____

This explains why large language models ...

- invent things and cannot cite sources;
- never produce the same answers;
- cannot count, compute, or reason*;
- can hardly correct their own mistakes once they have been made.

*: At least not with a vanilla transformer and a greedy decoding strategy.

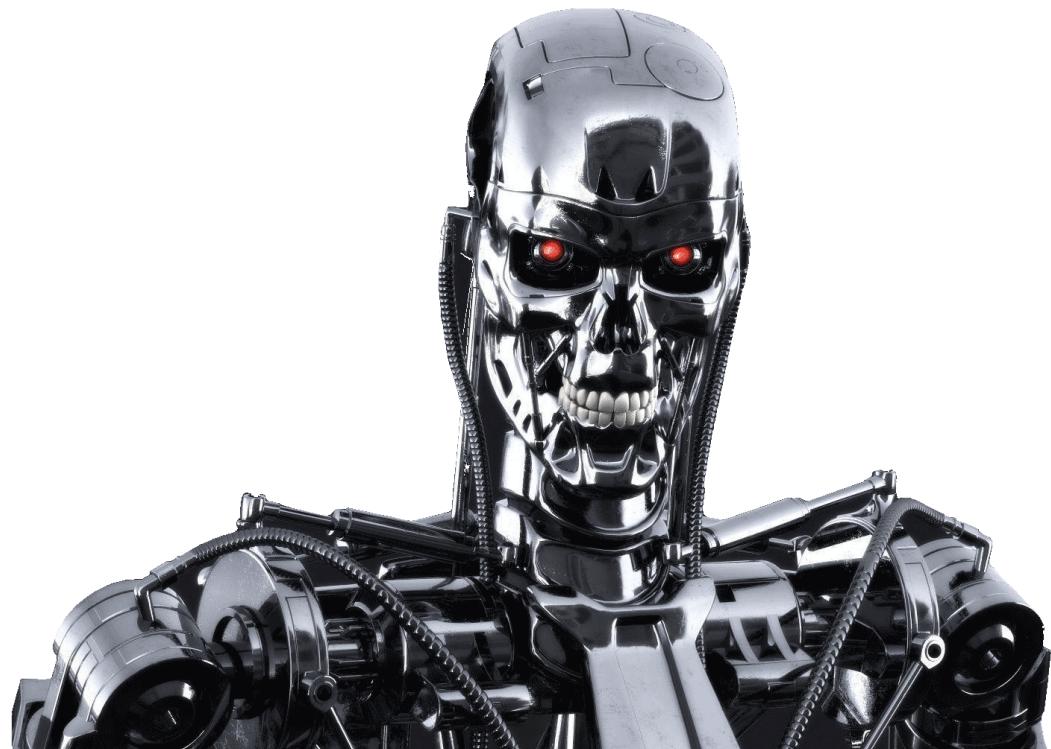


Rock, Paper, Scissors with GPT-4o



Not just text, but also images and sounds.

Artificial Intelligence



"With artificial intelligence we are summoning the demon" -- Elon Musk, 2014.



"We're really closer to a smart washing machine than Terminator" -- Fei-Fei Li,
Director of Stanford AI Lab, 2017.



AI 'godfather' quits Google over dangers of ...



Later bekij...



Delen

'AI could
be smarter
than us'

BBC NEWS



Geoffrey Hinton, 2023.

IMAGINATION
IN ACTION

Yann LeCun | Imagination In Action | Davos ...



Later bekij...
[Later bekijken](#)

Delen
[Delen](#)



Yann LeCun, 2023.


A definition of AI?



"Artificial intelligence is the science of making machines do things that would require intelligence if done by men." -- Marvin Minsky, 1968.

The Turing test

A computer passes the **Turing test** (aka the Imitation Game) if a human operator, after posing some written questions, cannot tell whether the written responses come from a person or from a computer.



*Can machines think?
(Alan Turing, 1950)*

An agent would not pass the Turing test without the following **requirements**:

- natural language processing
- knowledge representation
- automated reasoning
- machine learning
- computer vision (total Turing test)
- robotics (total Turing test)

Despite being proposed almost 70 years ago, the Turing test is **still relevant** today.

The Turing test tends to focus on **human-like errors**, **linguistic tricks**, etc.

However, it seems more important to study the **principles** underlying intelligence than to replicate an exemplar.



Aeronautics is not defined as the field of making machines
that fly so exactly like pigeons that they can fool even other pigeons.

A modern definition of AI

An 'AI system' is a machine-based system that is designed to operate with varying levels of autonomy and that may exhibit adaptiveness after deployment, and that, for explicit or implicit objectives, infers, from the input it receives, how to generate outputs such as predictions, content, recommendations, or decisions that can influence physical or virtual environments. -- European AI Act, Article 3, 2024.

A short history of AI

1940-1950: Early days

- 1943: McCulloch and Pitts: Boolean circuit model of the brain.
- 1950: Turing's "Computing machinery and intelligence".

1950-1970: Excitement and expectations

- 1950s: Early AI programs, including Samuel's checkers program, Newell and Simon's Logic Theorist and Gelernter's Geometry Engine.
- 1956: Dartmouth meeting: "Artificial Intelligence" adopted.
- 1958: Rosenblatt invents the perceptron.
- 1965: Robinson's complete algorithm for logical reasoning.
- 1966-1974: AI discovers computational complexity.



The Dartmouth workshop (1956)

The study is to proceed on the basis of the conjecture that every aspect of learning or any other feature of intelligence can in principle be so precisely described that a machine can be made to simulate it.



The Thinking Machine (Artificial Intelligenc...



Later bekij...



Delen



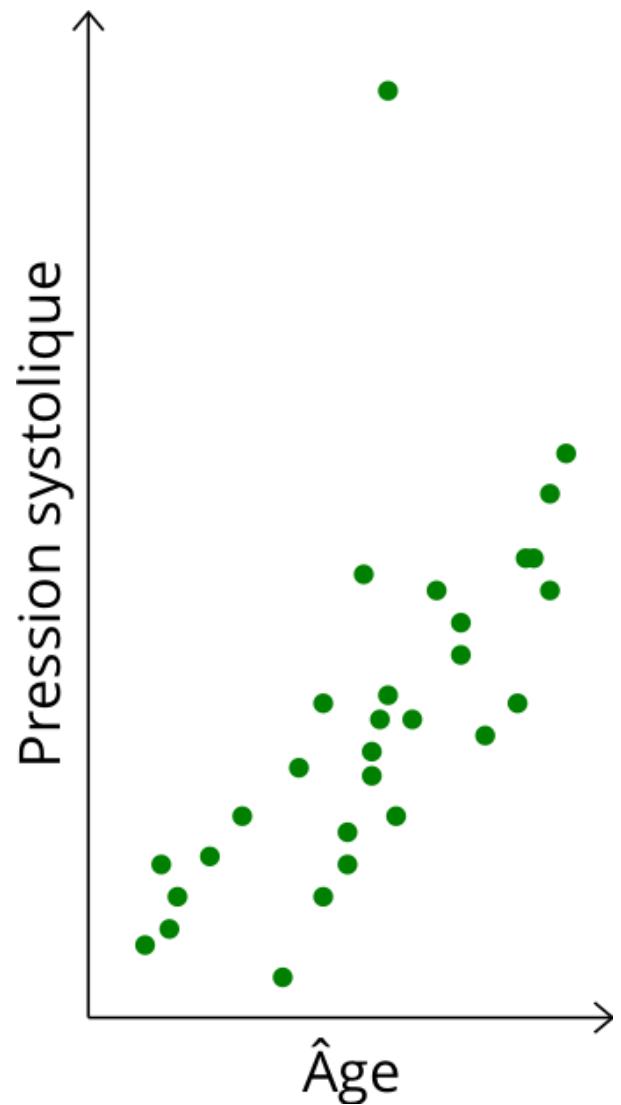
1970-1990: Knowledge-based approaches

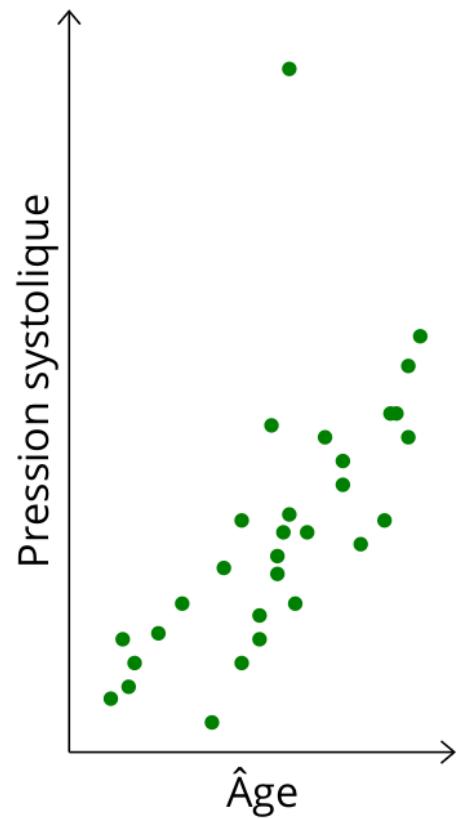
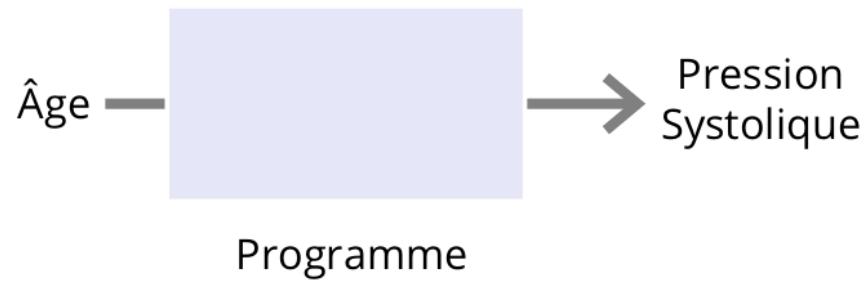
- 1969: Neural network research almost disappears after Minsky and Papert's book (1st AI winter).
- 1969-1979: Early development of knowledge-based systems.
- 1980-1988: Expert systems industrial boom.
- 1988-1993: Expert systems industry busts (2nd AI winter).

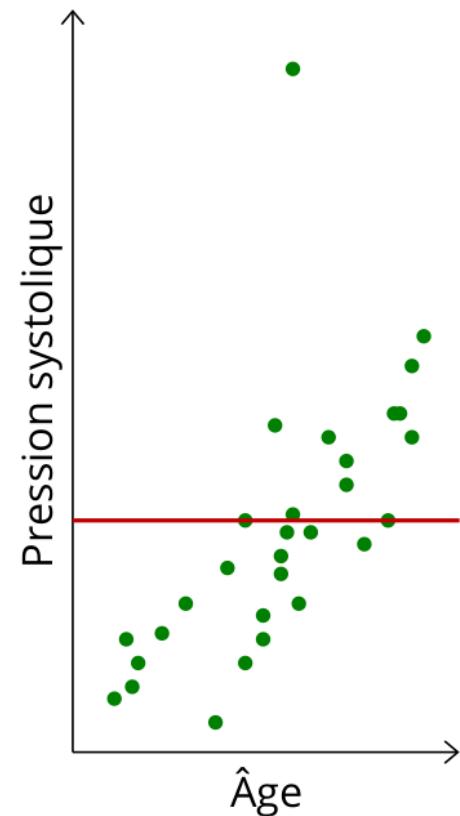
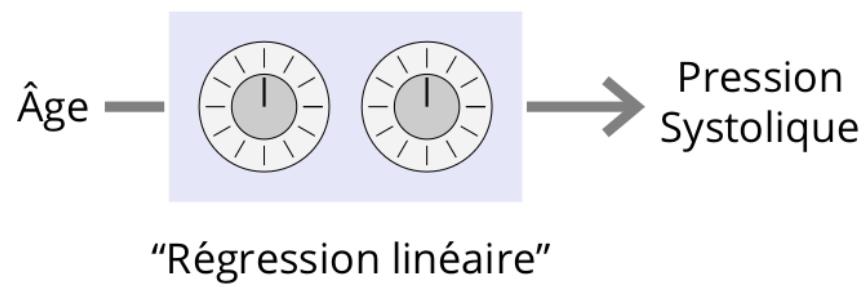
1990-Present: Statistical approaches

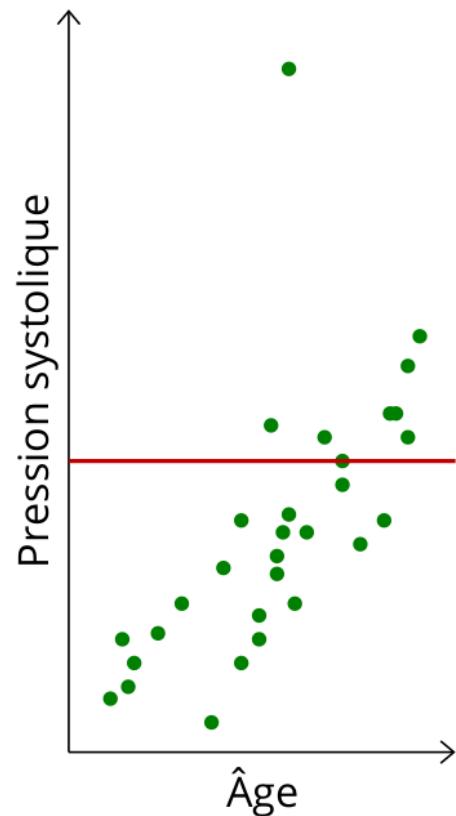
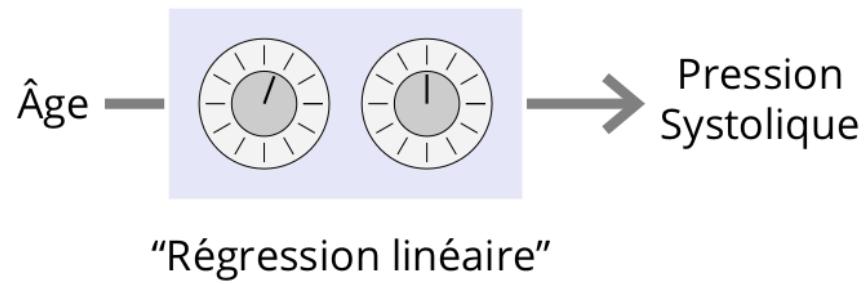
- 1985-1995: The return of neural networks.
- 1988-: Resurgence of probability, focus on uncertainty, general increase in technical depth.
- 1995-2010: New fade of neural networks.
- 2000-: Availability of very large datasets.
- 2010-: Availability of fast commodity hardware (GPUs).
- 2012-: Resurgence of neural networks with deep learning approaches.
- 2017: Attention is all you need (transformers).
- 2022: ChatGPT released to the public.

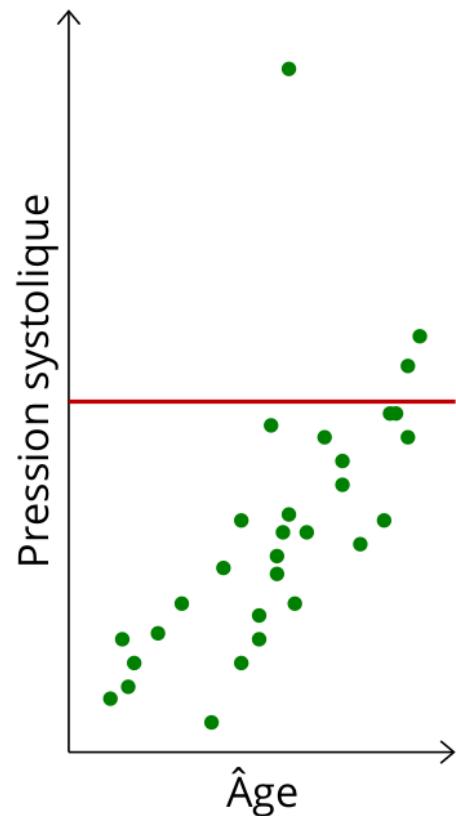
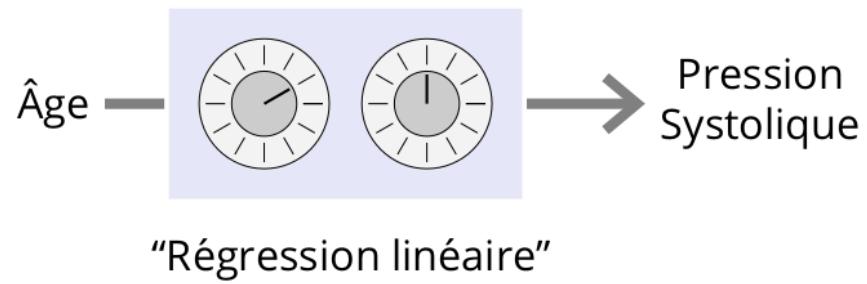
The deep learning revolution


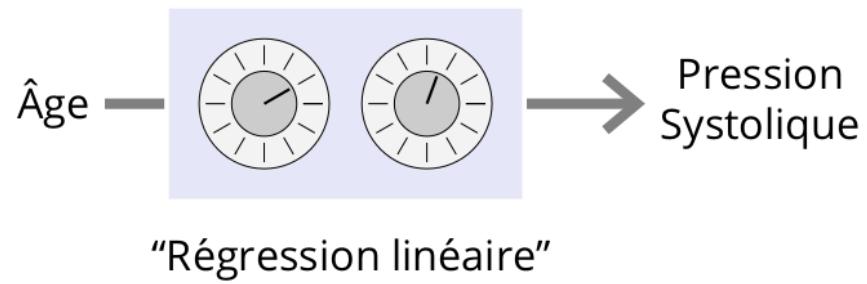


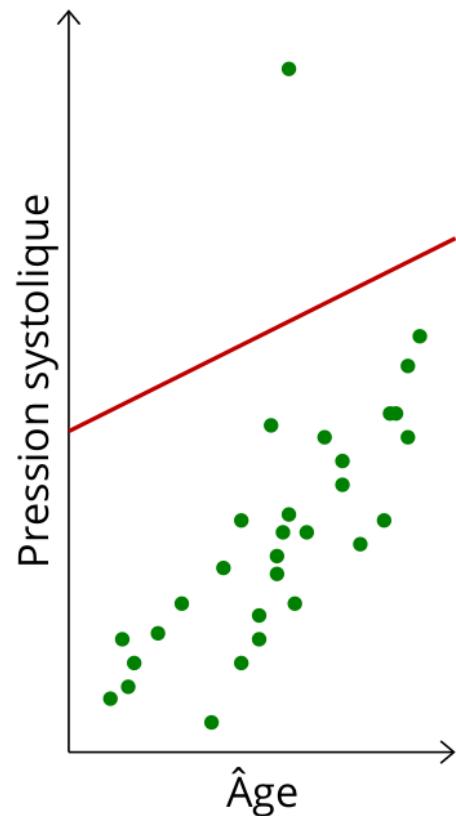
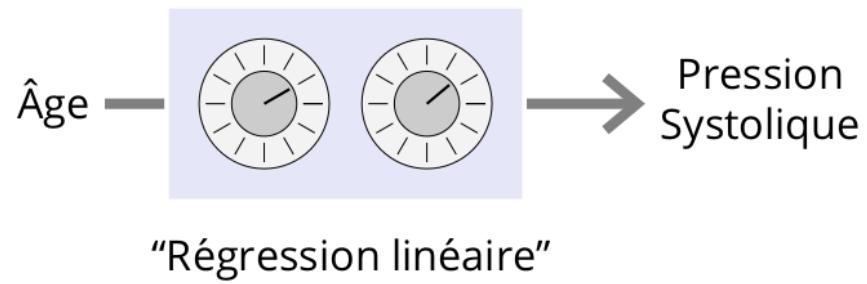


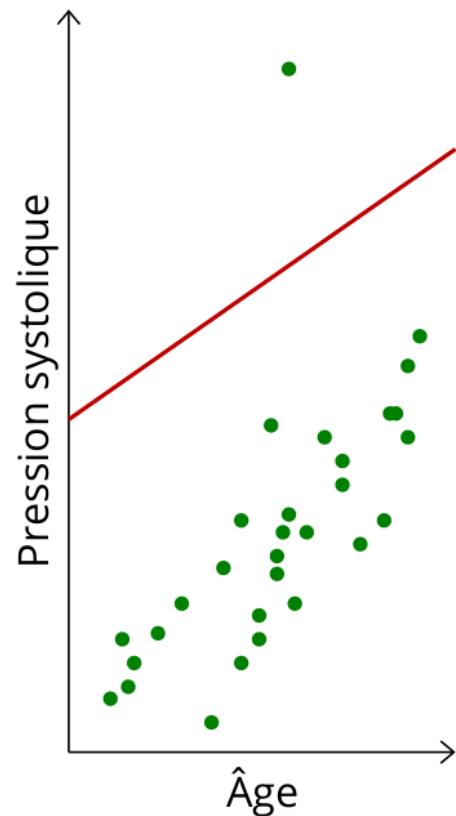
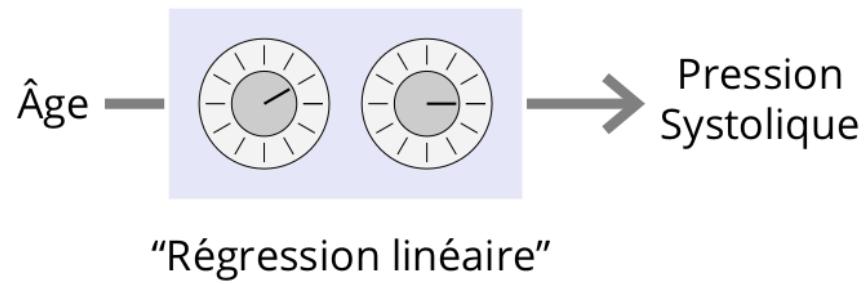


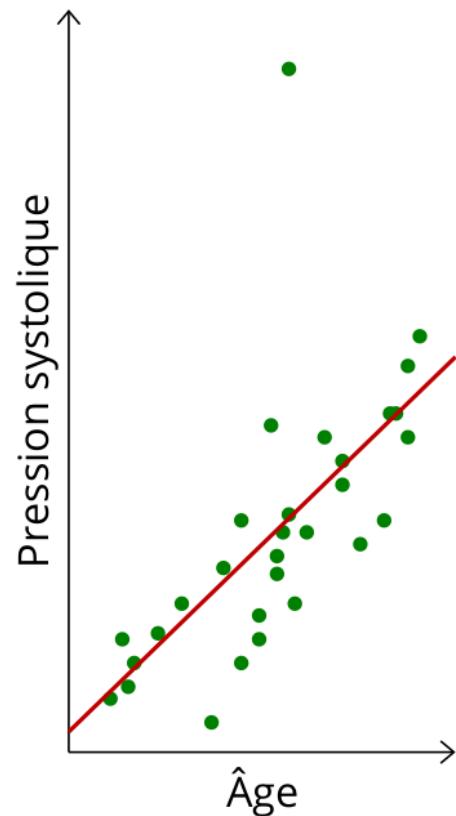
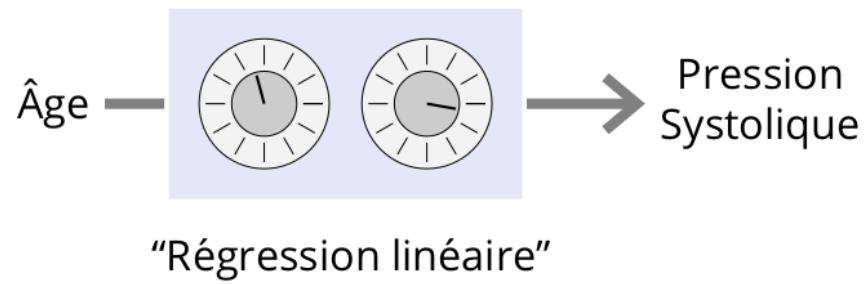






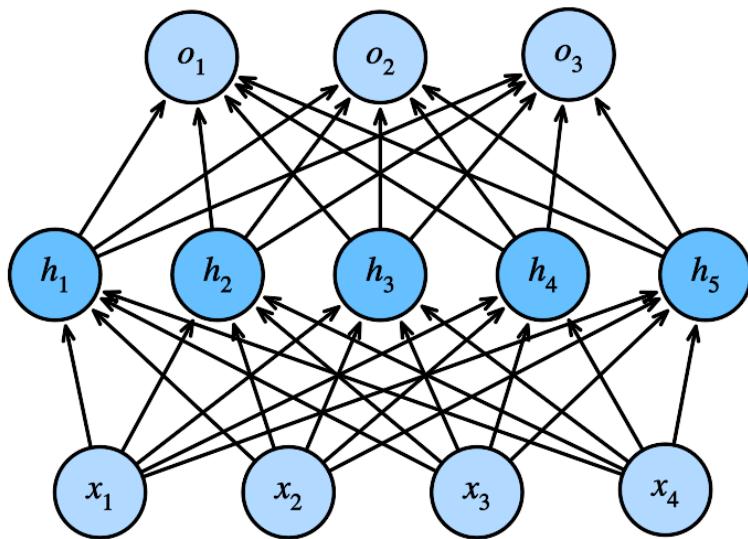




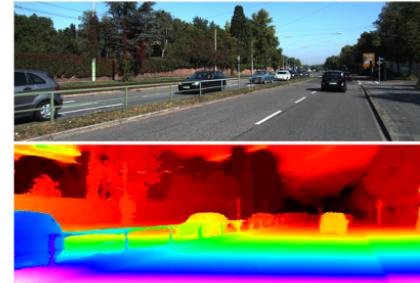
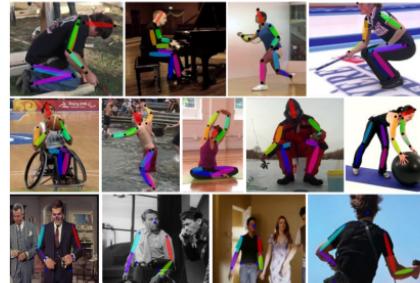
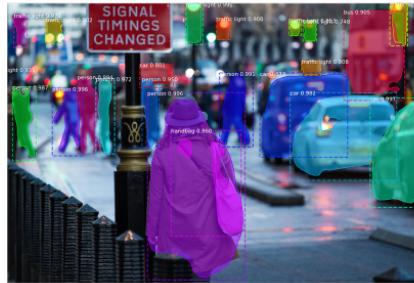


Deep learning **scales up** the statistical and machine learning approaches by

- using larger models known as neural networks,
- training on larger datasets,
- using more compute resources.



Specialized neural networks can be trained achieve super-human performance on many complex tasks that were previously thought to be out of reach for machines.

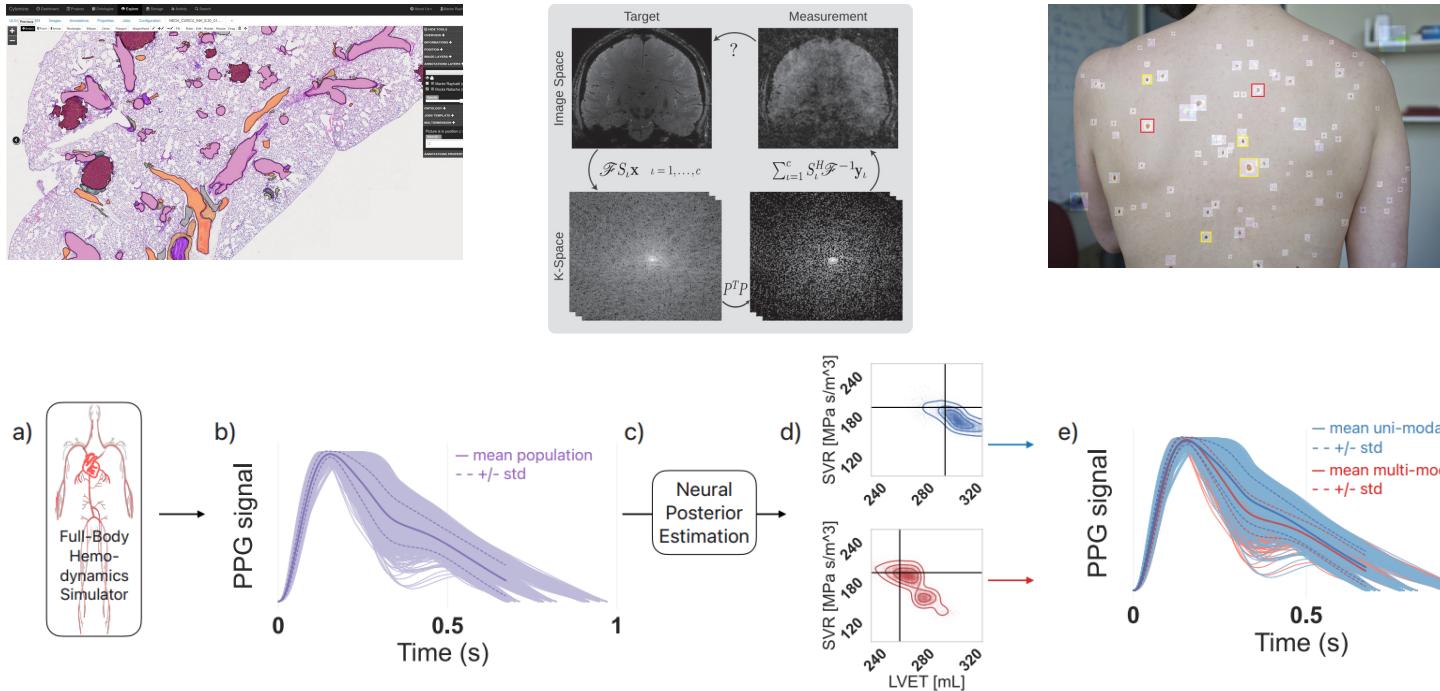


- I: Jane went to the hallway.
I: Mary walked to the bathroom.
I: Sandra went to the garden.
I: Daniel went back to the garden.
I: Sandra took the milk there.
Q: Where is the milk?
A: garden

(Top) Scene understanding, pose estimation, geometric reasoning.

(Bottom) Planning, Image captioning, Question answering.

Neural networks form **primitives** that can be transferred to many domains.



(Top) Analysis of histological slides, denoising of MRI images, nevus detection.

(Bottom) Whole-body hemodynamics reconstruction from PPG signals.

The breakthrough

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu


Lukasz Kaiser*
Google Brain
lukasz.kaiser@google.com

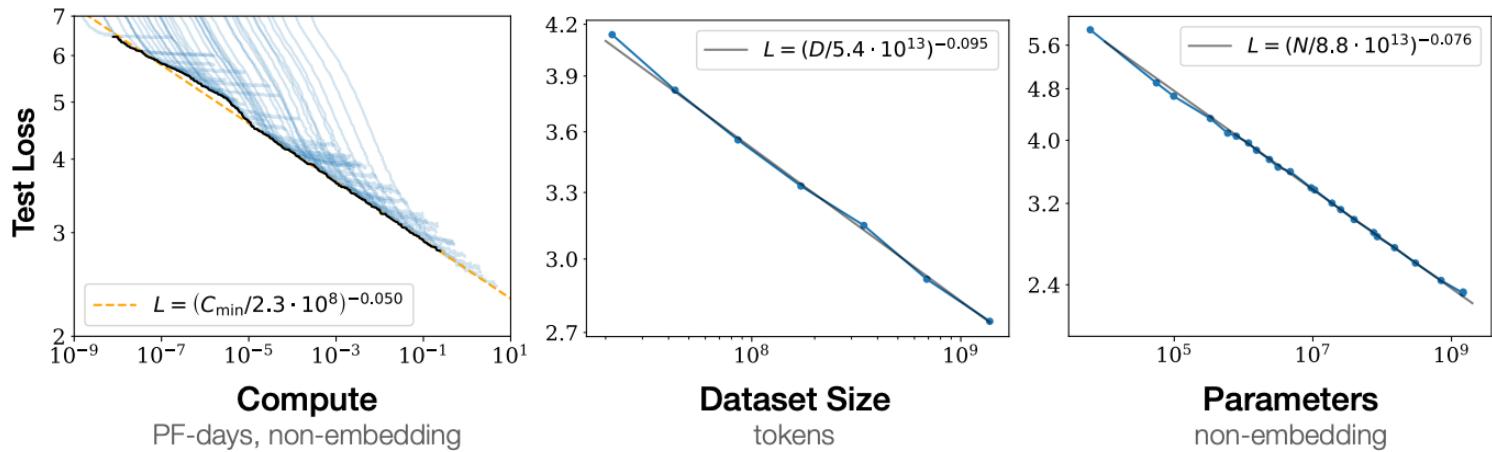
Illia Polosukhin* ‡
illia.polosukhin@gmail.com

Abstract

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks that include an encoder and a decoder. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely. Experiments on two machine translation tasks show these models to be superior in quality while being more parallelizable and requiring significantly less time to train. Our model achieves 28.4 BLEU on the WMT 2014 English-to-German translation task, improving over the existing best results, including ensembles, by over 2 BLEU. On the WMT 2014 English-to-French translation task, our model establishes a new single-model state-of-the-art BLEU score of 41.8 after training for 3.5 days on eight GPUs, a small fraction of the training costs of the best models from the literature. We show that the Transformer generalizes well to other tasks by applying it successfully to English constituency parsing both with large and limited training data.

Vaswani et al., 2017.





A brutal simplicity:

- The more data, the better the model.
- The more parameters, the better the model.
- The more compute, the better the model.

Scaling up further to gigantic models, datasets, and compute resources keeps pushing the boundaries of what is possible, **with no sign of slowing down**.



Claude 3.5 Sonnet as a writing partner



Later bekij...



Delen

The screenshot shows a conversational interface with two main panels. The left panel is a message thread where a user asks about a story involving Claw'd the Sarcastic Crab's Optimistic Mystery. The AI responds by suggesting to explore a sample dialogue between Claw'd and Dr. Wave. The right panel displays a 'Sample Dialogue: Claw'd and Dr. Wave's Translated Interaction' document. This document includes a setting description, a transcript of their interaction, and a translated output. A large red play button is overlaid on the right panel.

AI: +
Claw'd the Sarcastic Crab's Optimistic Mystery

Would you like me to explain any part of the genogram in more detail or explore how these family connections might influence the story?

5 Fantastic. I'm curious how best to showcase the arthropod <-> human interaction. Could you ideate on a sample of the translation?

Certainly! Let's explore a sample of how the arthropod-human interaction and translation might work between Claw'd and Dr. Wave. I'll create a brief dialogue that demonstrates this unique communication system.

Sample Dialogue: Claw'd and Dr. Wave's Translated Interaction
Click to open document

Claw'd: [Performs a series of rapid claw clicks, antennae vibrations, and releases a pattern of small bubbles]

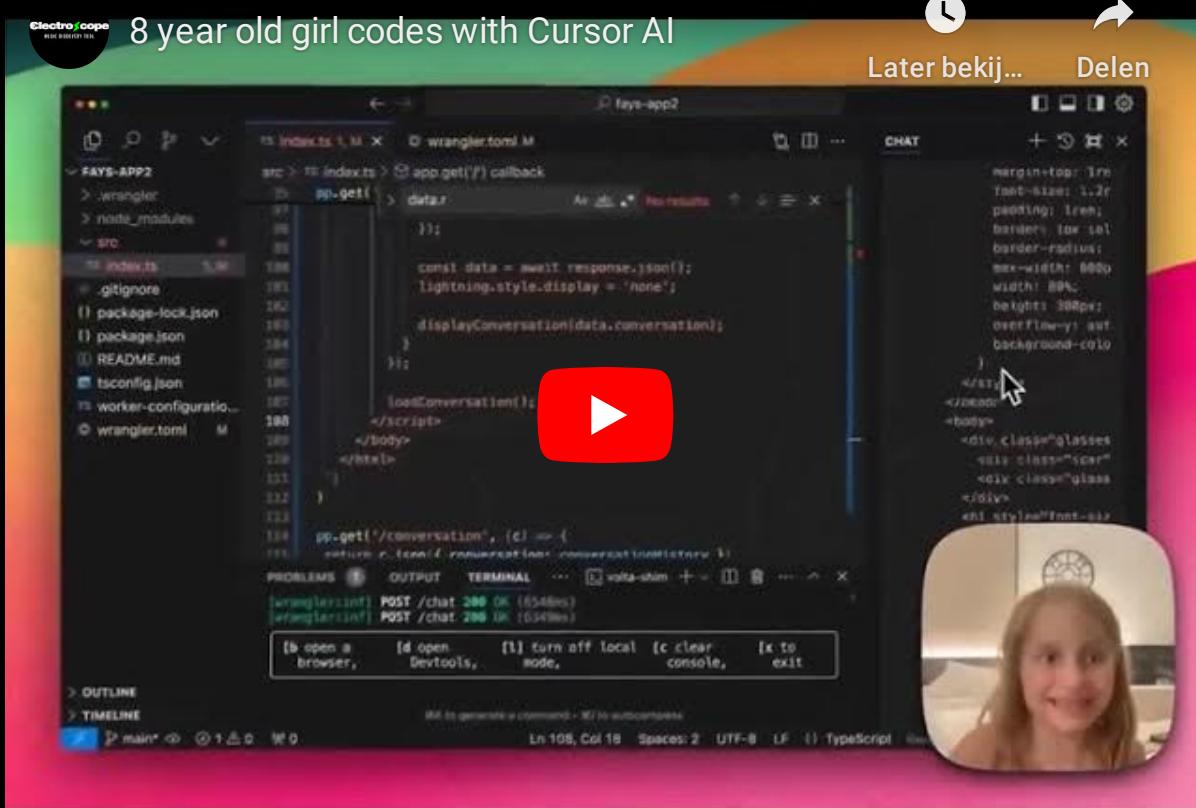
JustAI (Translation Interface): Analyzing crab communication patterns... Processing...

Translated output: "Hey doc,

Last edited just now

Copy

Conversational AI assistants (Anthropic, 2024)



Code assistants (Cursor, 2024)



Sense, Solve, and Go: The Magic of the Wa...



Later bekij...



Delen



Autonomous cars (Waymo, 2022)



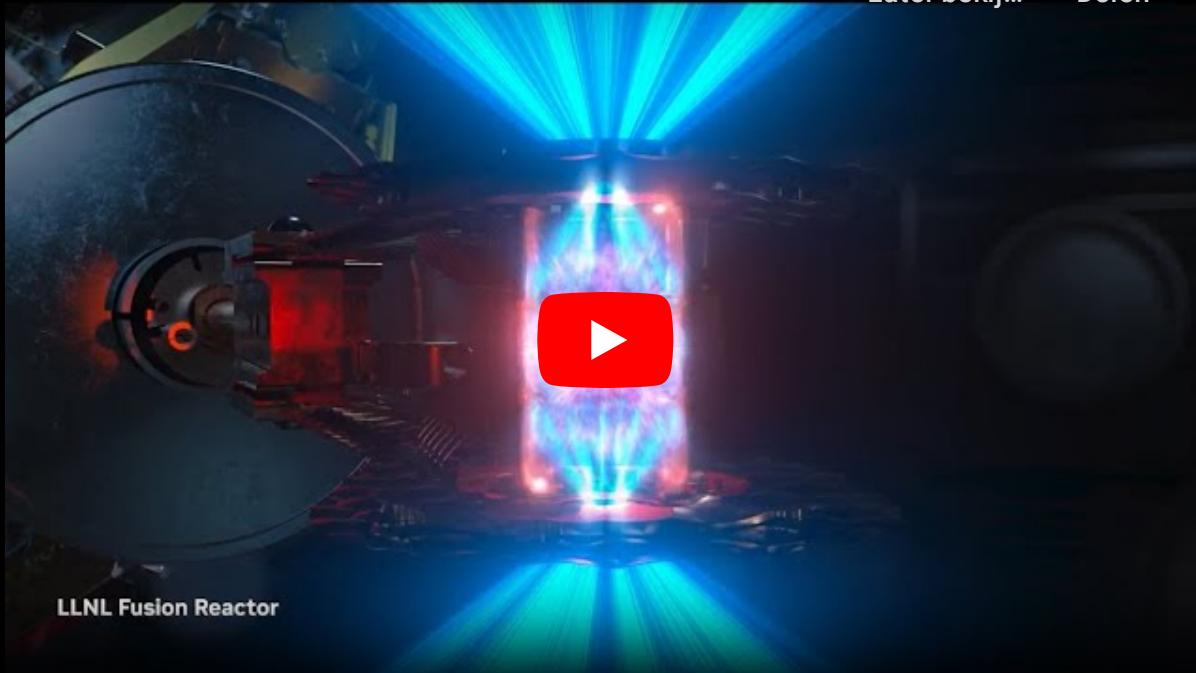
Powering the Future of Clean Energy | I AM ...



Later bekij...
Later bekijken



Delen



Powering the future of clean energy (NVIDIA, 2023)



Camels, Code & Lab Coats: How AI Is Advancing Medicine



Later bekijk...



Delen



How AI is advancing medicine (Google, 2018)

Deep learning can also **solve problems that no one could solve before**.

AlphaFold: From a sequence of amino acids to a 3D structure

nature

Explore content ▾ About the journal ▾ Publish with us ▾

nature > articles > article

Article | [Open access](#) | Published: 15 July 2021

Highly accurate protein structure prediction with AlphaFold

John Jumper , Richard Evans, Alexander Pritzel, Tim Green, Michael Figurnov, Olaf Ronneberger, Kathryn Tunyasuvunakool, Russ Bates, Augustin Žídek, Anna Potapenko, Alex Bridgland, Clemens Meyer, Simon A. Kohl, Andrew J. Ballard, Andrew Cowie, Bernardino Romera-Paredes, Stanislav Nikolov, Rishabh Jain, Jonas Adler, Trevor Back, Stig Petersen, David Reiman, Ellen Clancy, Michał Zieliński, ... Demis Hassabis 

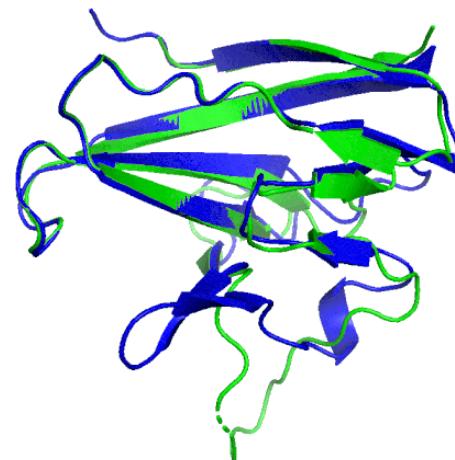
+ Show authors

[Nature](#) 596, 583–589 (2021) | [Cite this article](#)

1.42m Accesses | 12k Citations | 3493 Altmetric | [Metrics](#)

Abstract

Proteins are essential to life, and understanding their structure can facilitate a mechanistic understanding of their function. Through an enormous experimental effort^{1,2,3,4}, the structures of around 100,000 unique proteins have been determined⁵, but this represents a small fraction of the billions of known protein sequences^{6,7}. Structural coverage is bottlenecked by the months to years of painstaking effort required to determine a single protein structure. Accurate computational approaches are needed to address this gap and to enable large-scale structural bioinformatics. Predicting the three-dimensional structure that a protein will adopt based solely on its amino acid sequence—the structure prediction component of the ‘protein folding problem’⁸—has been an important open research problem for more than 50 years⁹. Despite recent progress^{10,11,12,13,14}, existing methods fall far short of atomic accuracy, especially when no homologous structure is available. Here we provide the






AlphaFold: The making of a scientific break...



Later bekij...



Delen



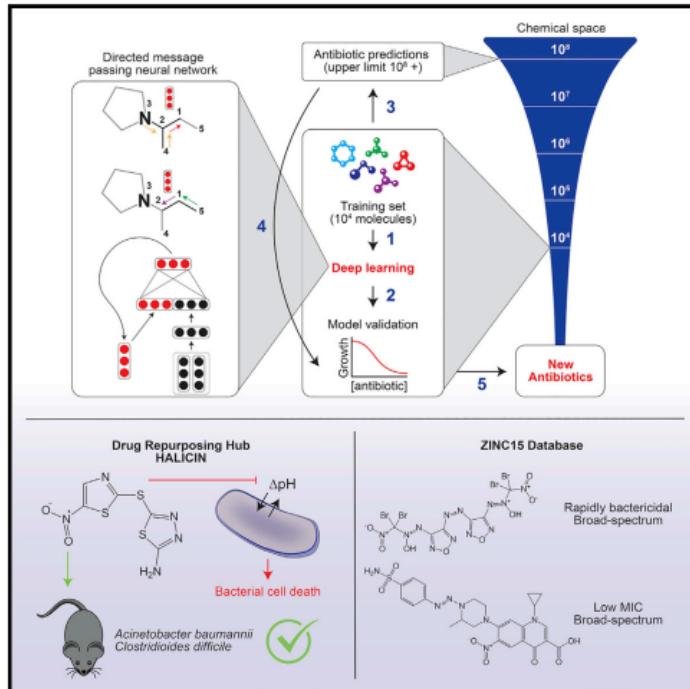
AI for Science (Deepmind, AlphaFold, 2020)

Drug discovery with graph neural networks

Cell

A Deep Learning Approach to Antibiotic Discovery

Graphical Abstract



Authors

Jonathan M. Stokes, Kevin Yang,
Kyle Swanson, ..., Tommi S. Jaakkola,
Regina Barzilay, James J. Collins

Correspondence


regina@csail.mit.edu (R.B.),
jimjc@mit.edu (J.J.C.)

In Brief


A trained deep neural network predicts antibiotic activity in molecules that are structurally different from known antibiotics, among which Halicin exhibits efficacy against broad-spectrum bacterial infections in mice.

GraphCast: fast and accurate weather forecasts


a) Input weather state




b) Predict the next state




c) Roll out a forecast




d) Encoder










e) Processor



f) Decoder




g) Simultaneous multi-mesh message-passing



INFO8006 Introduction to AI

Course outline

- Lecture 0: Artificial intelligence
- Lecture 1: Intelligent agents
- Lecture 2: Solving problems by searching
- Lecture 3: Adversarial search
- Lecture 4: Quantifying uncertainty
- Lecture 5: Probabilistic reasoning
- Lecture 6: Reasoning over time
- Lecture 7: Machine learning and neural networks
- Lecture 8: Making decisions
- Lecture 9: Reinforcement learning



My mission

By the end of this course, you will have built autonomous agents that efficiently make decisions in fully informed, partially observable and adversarial settings. Your agents will draw inferences in uncertain and unknown environments and optimize actions for arbitrary reward structures.

The models and algorithms you will learn in this course apply to a wide variety of artificial intelligence problems and will serve as the foundation for further study in any application area (from engineering and science, to business and medicine) you choose to pursue.

Goals and philosophy

General

- Understand the landscape of artificial intelligence.
- Be able to write from scratch, debug and run (some) AI algorithms.

Well-established and state-of-the-art algorithms

- Good old-fashioned AI: well-established algorithms for intelligent agents and their mathematical foundations.
- Introduction to materials new from research (\leq 5 years old).
- Understand some of the open questions and challenges in the field.

Practical

- Fun and challenging course projects.

